

FEATURES-DEMO

June 22, 2020

1 Stops, places and moves location analysis

Definitions: - **Location data** is collected as a sequence of location samples with varying sample frequency and accuracy. - **Places** are locations of relevance to the user, such as home or workplace and are described by their coordinates and an ID. - **Stops** are specific visits to one of those places, described by their coordinates along with arrival and departure time. A stop is always associated with exactly one place while a place can be associated with many stops. Stops are always non-overlapping in time. - **Moves** are sequences of location points between stops and are described by departure and arrival time, origin and destination place and the distance of the move.

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from geopy.distance import geodesic
from sklearn.cluster import DBSCAN
import gmaps

import sys
sys.path.append("../")
from location import *

# Keep data consistent, load from disk.
LOAD_DATA_FROM_DISK = False
```

```
In [2]: def get_date(row):
        return row.date()
```

1.1 Generate example data

```
In [3]: if LOAD_DATA_FROM_DISK:
        df = pd.read_json('multi_date_data.json').T
        df.datetime = pd.to_datetime(df.datetime, unit='ms')
        df.date = df.datetime.dt.date.astype('datetime64[ns]')
    else:
        # create data simulating 3 places (a,b,c)
        a = (55.686381, 12.557155) # Blaagaards Plads
        b = (55.666919, 12.536792) # Spaces
        c = (55.688305, 12.561862) # Hulen
```

```

X = np.vstack([
    # day 1: home, work, home, workout, home
    np.array([a]*(60*8+30)),
    np.array([np.linspace(a[0], b[0], 30), np.linspace(a[1], b[1], 30)]).T,
    np.array([b]*(60*7+30)),
    np.array([np.linspace(b[0], a[0], 30), np.linspace(b[1], a[1], 30)]).T,
    np.array([a]*55),
    np.array([np.linspace(a[0], c[0], 5), np.linspace(a[1], c[1], 5)]).T,
    np.array([c]*55),
    np.array([np.linspace(c[0], a[0], 5), np.linspace(c[1], a[1], 5)]).T,
    np.array([a]*60*5),
    # day 2: home, work, home
    np.array([a]*(60*8+30)),
    np.array([np.linspace(a[0], b[0], 30), np.linspace(a[1], b[1], 30)]).T,
    np.array([b]*(60*7+30)),
    np.array([np.linspace(b[0], a[0], 30), np.linspace(b[1], a[1], 30)]).T,
    np.array([a]*60*7),
    # day 3: home, workout, home
    np.array([a]*(60*10+55)),
    np.array([np.linspace(a[0], c[0], 5), np.linspace(a[1], c[1], 5)]).T,
    np.array([c]*55),
    np.array([np.linspace(c[0], a[0], 5), np.linspace(c[1], a[1], 5)]).T,
    np.array([a]*60*12),
])

X += np.random.normal(loc=0, scale=0.00005, size=X.shape)

df = pd.DataFrame(X, columns=['latitude', 'longitude'])
df.insert(0, 'user_id', 0)
df.insert(1, 'timestamp', np.arange(df.shape[0]) * 60000 + 1573430400000.0)
df.insert(2, 'datetime', pd.to_datetime(df.timestamp, unit='ms'))
df.insert(3, 'date', df.datetime.dt.date.astype('datetime64[ns]'))

# Write to file
df.T.to_json('multi_date_data.json')

df.head()

df.date = df.date.apply(get_date)


In [57]: dates = np.unique(df.date.values)
         dates

Out[57]: array([datetime.date(2019, 11, 11), datetime.date(2019, 11, 12),
               datetime.date(2019, 11, 13)], dtype=object)

```

1.2 Quick visualization of the 3 places visited

```
In [58]: plt.figure(figsize=(8,8))
plt.plot(df.latitude.values, df.longitude.values, marker='.', alpha=1)
plt.grid()
plt.show()
```



1.3 Preprocessing (stops, places and moves)

- A stop is a collection of stationary points
- A place is a cluster of stops found using DBSCAN
- A move is a transition from one stop to another.

```
In [6]: stops, places, moves = get_stops_places_and_moves(df)
stops['date'] = stops.arrival.apply(get_date)
moves['date'] = moves.arrival.apply(get_date)
```

In [7]: stops

```
Out[7]:
```

	user_id	latitude	longitude	samples	arrival \
0	0	55.686381	12.557161	511	2019-11-11 00:00:00
1	0	55.666914	12.536788	452	2019-11-11 08:59:00
2	0	55.686383	12.557170	57	2019-11-11 16:59:00
3	0	55.688303	12.561877	57	2019-11-11 17:59:00
4	0	55.686376	12.557154	812	2019-11-11 18:59:00
5	0	55.666921	12.536791	452	2019-11-12 08:59:00
6	0	55.686382	12.557154	1077	2019-11-12 16:59:00
7	0	55.688310	12.561865	57	2019-11-13 10:59:00
8	0	55.686377	12.557154	721	2019-11-13 11:59:00

	departure	duration	place	date
0	2019-11-11 08:30:00	510.0	0	2019-11-11
1	2019-11-11 16:30:00	451.0	1	2019-11-11
2	2019-11-11 17:55:00	56.0	0	2019-11-11
3	2019-11-11 18:55:00	56.0	2	2019-11-11
4	2019-11-12 08:30:00	811.0	0	2019-11-11
5	2019-11-12 16:30:00	451.0	1	2019-11-12
6	2019-11-13 10:55:00	1076.0	0	2019-11-12
7	2019-11-13 11:55:00	56.0	2	2019-11-13
8	2019-11-13 23:59:00	720.0	0	2019-11-13

In [8]: places

```
Out[8]:
```

	user_id	place	latitude	longitude	duration	stops
0	0	0	55.686381	12.557154	3173.0	5
1	0	1	55.666918	12.536790	902.0	2
2	0	2	55.688307	12.561871	112.0	2

In [9]: moves

```
Out[9]:
```


	user_id	from_latitude	from_longitude	to_latitude	to_longitude	samples \
0	0	55.686350	12.557288	55.666931	12.536812	30
1	0	55.666994	12.536835	55.686413	12.557147	30
2	0	55.686326	12.557210	55.666843	12.536827	30
3	0	55.666941	12.536796	55.686385	12.557099	30

	departure	arrival	from_place	to_place	distance \
0	2019-11-11 08:30:00	2019-11-11 08:59:00	0.0	1.0	2520.430627
1	2019-11-11 16:30:00	2019-11-11 16:59:00	1.0	0.0	2516.562350
2	2019-11-12 08:30:00	2019-11-12 08:59:00	0.0	1.0	2523.909200
3	2019-11-12 16:30:00	2019-11-12 16:59:00	1.0	0.0	2518.612156

	duration	mean_speed	date
0	29.0	1.448523	2019-11-11
1	29.0	1.446300	2019-11-11
2	29.0	1.450523	2019-11-12
3	29.0	1.447478	2019-11-12

1.3.1 Visualizing the clusters and moves

```
In [59]: plt.figure(figsize=(8,8))
plt.plot(df.longitude.values, df.latitude.values, marker='.', alpha=.2)
plt.scatter(stops.longitude.values, stops.latitude.values, marker='d', color='r', zorder=2)
plt.scatter(places.longitude.values, places.latitude.values, s=1000, facecolors='none', zorder=3)
for index, move in moves.iterrows():
    plt.plot([move.from_longitude, move.to_longitude], [move.from_latitude, move.to_latitude], color='r', zorder=1)
plt.grid()
plt.show()
```



2 Features

2.1 Number of clusters

This feature represents the total number of clusters found by the clustering algorithm.

```
In [60]: def number_of_clusters(places):
         return len(places)
```

```
In [61]: number_of_clusters(places)
```

```
Out[61]: 3
```

2.2 Location Variance:

This feature measures the variability of a participant's location data from stationary states. LV was computed as the natural logarithm of the sum of the statistical variances of the latitude and the longitude components of the location data.

```
In [62]: def location_variance(df):
         # If fewer than 2 observations, we can't compute the variance
         if len(df) < 2:
             return 0.0
         return np.log(df.latitude.var() + df.longitude.var() + 1)
```

```
In [63]: location_variance(df)
```

```
Out[63]: 0.00013597465949774686
```

2.3 Location Entropy (LE):

A measure of points of interest. High entropy indicates that the participant spent time more uniformly across different location clusters, while lower entropy indicates the participant spent most of the time at some specific clusters. Concretely it is calculated as:

$$Entropy = - \sum_{i=1}^N p_i \cdot \log p_i$$

where each i represents a location cluster, N denotes the total number of location clusters, and p_i is the percentage of time the participant spent at the location cluster i . High cluster entropy indicates that the participant spent time more uniformly across different location clusters, while lower cluster entropy indicates the participant spent most of the time at some specific clusters.

Here, we use the duration spent at each place, found in the duration column in the places dataframe.

```
In [64]: def _entropy(durations):
         p = durations / np.sum(durations)
         return -np.sum(p * np.log(p))
```

```
In [65]: _entropy(places.duration)
```

```
Out[65]: 0.6377255748619863
```

```
In [66]: # NumPy for reference:
```

```
from scipy.stats import entropy
entropy(places.duration)
```

```
Out[66]: 0.6377255748619863
```

2.4 Normalized LE:

Normalized entropy is calculated by dividing the cluster entropy by its maximum value, which is the logarithm of the total number of clusters. Normalized entropy is invariant to the number of clusters and thus solely depends on their visiting distribution. The value of normalized entropy ranges from 0 to 1, where 0 indicates the participant has spent their time at only one location, and 1 indicates that the participant has spent an equal amount of time to visit each location cluster.

Here we just divide by the log to the number of places.

```
In [67]: def normalized_entropy(durations):  
         return entropy(durations) / np.log(len(durations))
```

```
In [68]: normalized_entropy(places.duration)
```

```
Out[68]: 0.5804828340625297
```

2.5 Transition Time:

Transition Time measures the percentage of time the participant has been in the transition state.

A few ways of doing this, but one is using the moves dataframe and simply summing the duration column, and dividing by 24 hours.

```
In [69]: def transition_time(moves):  
         move_time = moves.duration.sum()  
         return move_time / (24 * 60)
```

```
In [70]: transition_time(moves)
```

```
Out[70]: 0.08055555555555556
```

3 Total Distance:

This feature measures the total distance the participant has traveled in the transition state.

Here we simply sum the distance column in the moves dataframe.

```
In [71]: def total_distance(moves):  
         return moves.distance.sum()
```

```
In [72]: total_distance(moves)
```

```
Out[72]: 10079.514332342535
```

3.1 Routine Index

```
In [73]: HOURS_IN_A_DAY = 24
```

```
def print_hour_matrix(M):  
    for i, row in enumerate(M):  
        line = "[{:0>2} - {:0>2}] ".format(i, i+1)  
        for e in row:
```

```

        line += '%0.2f ' % e
    print(line)

def make_hour_matrix(stops, num_places):
    h = np.zeros((HOURS_IN_A_DAY, num_places))

    for index, row in stops.iterrows():
        pid = row.place
        start_hour = row.arrival.hour
        end_hour = row.departure.hour

        # If user arrived and departed within the same hour
        # Then the time stayed is the diff between departure and arrival
        if start_hour == end_hour:
            h[start_hour, pid] = row.departure.minute - row.arrival.minute

        else:
            # Arrival hour
            h[start_hour, pid] = 60 - row.arrival.minute

            # In between
            for hour in range(start_hour+1, end_hour):
                h[hour, pid] = 60

            # Departure hour
            h[end_hour, pid] = row.departure.minute

    return h / 60 # Normalize by 60 mins

In [74]: # Plot a matrix as a color map
def matrix_plot(m):
    plt.figure(figsize=(10,10))
    plt.imshow(m, cmap='bone')
    plt.title('Hour matrix')
    plt.xlabel('Place ID')
    plt.ylabel('Timeslot')
    plt.yticks(range(HOURS_IN_A_DAY), ["{:0>2} - {:0>2} ".format(i, i+1) for i in range(24)])
    plt.xticks(range(m.shape[1]))
    plt.show()

In [75]: s1 = stops[stops.date == dates[0]]
s1


Out[75]:
```

	user_id	latitude	longitude	samples	arrival \
0	0	55.686381	12.557161	511	2019-11-11 00:00:00
1	0	55.666914	12.536788	452	2019-11-11 08:59:00
2	0	55.686383	12.557170	57	2019-11-11 16:59:00
3	0	55.688303	12.561877	57	2019-11-11 17:59:00


```
4          0  55.686376  12.557154      812 2019-11-11 18:59:00
```

	departure	duration	place	date
0	2019-11-11 08:30:00	510.0	0	2019-11-11
1	2019-11-11 16:30:00	451.0	1	2019-11-11
2	2019-11-11 17:55:00	56.0	0	2019-11-11
3	2019-11-11 18:55:00	56.0	2	2019-11-11
4	2019-11-12 08:30:00	811.0	0	2019-11-11

```
In [76]: h1 = make_hour_matrix(s1, len(places))  
         matrix_plot(h1)
```



```

In [77]: def RI(h_mean, h, end_hour=24):
        '''
        input:
            h_mean (2d matrix): Historical Mean Matrix
            h (2d matrix): Hour Matrix for a day

        output:
            routine_index: -1 (could not be calculated) or [0 to 1].
        '''
        if h_mean.sum() == 0:
            return -1.0 # no routine index could be calculated

        assert(h_mean.shape == h.shape)

        m,n = h.shape

        overlap = 0.0
        for i in range(m):
            for j in range(n):
                overlap += min(h_mean[i,j], h[i,j])

        max_overlap = min(h_mean.sum(), h.sum())

        return overlap / max_overlap

```

3.2 Using today's stops and historical stops to calculate routine index

I.e. no updating of routine matrix, always recalculate it.

```

In [78]: STOPS = {}
        for date in dates:
            print('Date:', date)
            # Select data by date
            data = df[df.date == date]

            # Find stops, moves, places
            S, P, M = get_stops_places_and_moves_daily(data, merge=False, move_duration=3)

            # Store them
            STOPS[date] = S

```

Date: 2019-11-11

Date: 2019-11-12

Date: 2019-11-13

```

In [79]: def plot_today_and_routine(today, routine, routine_after, save=False):
        interval_strings = ["[:0>2} - {:0>2}] ".format(i, i+1) for i in range(HOURS_IN_A

```

```

f, (ax1, ax2, ax3) = plt.subplots(1, 3)
f.set_size_inches((10,10))

ax1.imshow(routine, cmap='bone')
ax1.set_title('Routine')
ax1.set_xlabel('Place ID')
ax1.set_ylabel('Timeslot')
ax1.set_yticks(range(HOURS_IN_A_DAY))
ax1.set_yticklabels(interval_strings)
ax1.set_xticks(range(routine.shape[1]))

ax2.imshow(today, cmap='bone')
ax2.set_title('Today')
ax2.set_xlabel('Place ID')
ax2.set_ylabel('Timeslot')
ax2.set_yticks(range(HOURS_IN_A_DAY))
ax2.set_yticklabels(interval_strings)
ax2.set_xticks(range(today.shape[1]))

ax3.imshow(routine_after, cmap='bone')
ax3.set_title('Updated Routine')
ax3.set_xlabel('Place ID')
ax3.set_ylabel('Timeslot')
ax3.set_yticks(range(HOURS_IN_A_DAY))
ax3.set_yticklabels(interval_strings)
ax3.set_xticks(range(routine_after.shape[1]))

if save:
    plt.savefig('routine.png')
plt.show()

```

```
In [80]: DISTF = lambda a, b: geodesic(a, b).meters
```

```

def get_places_2(stops, dist=25, distf=DISTF):
    if stops.empty:
        stops['place'] = []
        places = pd.DataFrame(columns=['user_id', 'place', 'latitude', 'longitude', 'duration', 'samples'])
    else:
        points = stops[['latitude', 'longitude']].values
        dbs = DBSCAN(dist, min_samples=1, metric=distf).fit(points)
        stops['place'] = dbs.labels_
        places = stops.groupby('place').agg({
            'latitude': np.median,
            'longitude': np.median,
            'duration': np.sum,
            'samples': len,
        }).reset_index()

```

```

        places.rename(columns={'samples': 'stops'}, inplace=True)
        places.insert(0, 'user_id', stops.user_id.values[0])
    return stops, places

```

```

In [96]: for date in dates:
        # Calculate today's matrix
        stops_today = STOPS[date]
        stops_so_far = [STOPS[d] for d in dates[dates <= date]]
        stops_so_far = pd.concat(stops_so_far)
        stops_so_far = stops_so_far.sort_values(['arrival'])
        stops_so_far = stops_so_far.reset_index()
        stops_so_far, places_so_far = get_places_2(stops_so_far)

        number_of_places = len(places_so_far)

        hour_matrix_today = make_hour_matrix(stops_today, number_of_places)

        dates_hist = dates[dates < date]

        routine_matrix = hour_matrix_today

        if len(dates_hist) > 0:
            print(date)
            hour_matrices_hist = [make_hour_matrix(STOPS[date_hist], number_of_places) for date_hist in dates_hist]
            new_routine_matrix = np.mean(hour_matrices_hist, axis=0)
            ri = RI(new_routine_matrix, hour_matrix_today)
            plot_today_and_routine(hour_matrix_today, routine_matrix, new_routine_matrix)
            routine_matrix = new_routine_matrix

        else:
            ri = -1
        print('Routine index: %0.2f' % ri)


        hms[date] = hour_matrix_today
        rms[date] = routine_matrix

        print('-'*35)

```


Routine index: -1.00

2019-11-12



Routine index: 0.96

2019-11-13



Routine index: 0.69

3.3 Home Stay:

The percentage of time the participant has been at the cluster that represents home. We define the home cluster as the cluster, which is mostly visited during the period between 12 am and 6 am.

Implementation steps: * Identify home: Use the hours dataframe to determine the most visited cluster between 00 and 06 am. * Count percentage of time at home: Use the places dataframe to calculate the time distribution.

However - we need to fill out the hours dataframe with data between 00 and 06 first

```
In [33]: date = dates[0]
         # Calculate today's matrix
         stops_today = STOPS[date]
         stops_so_far = [STOPS[d] for d in dates[dates <= date]]
         stops_so_far = pd.concat(stops_so_far)
         stops_so_far = stops_so_far.sort_values(['arrival'])
         stops_so_far = stops_so_far.reset_index()
         stops_so_far, places_so_far = get_places_2(stops_so_far)

         num_places = len(places_so_far)

In [34]: H = make_hour_matrix(stops_today, num_places)

In [104]: def get_home_place(hour_matrix):
          start, end = 0, 6
          place_dist = hour_matrix[start:end].sum()

          # Check that there was actually data between 00 and 06
          assert not np.all(hour_matrix[start:end].sum() == 0)

          return hour_matrix[start:end].sum().argmax()

          def home_stay(places, hour_matrix):
              distr = places.duration / places.duration.sum()
              home_id = get_home_place(hour_matrix)
              return distr[home_id]

In [106]: get_home_place(H)

Out[106]: 0

In [109]: home_stay(places_so_far, H)

Out[109]: 0.6307356154406409
```